

TMA4268 Statistical Learning

Module 6: Solution sketches

Sara Martino, Stefanie Muff, Kenneth Aase, Daesoo Lee
Department of Mathematical Sciences, NTNU

February 22, 2024

Recommended exercise 1

For the least square estimator, the solution can be found in the first section here.

For the maximum likelihood estimator, the solution can be found here.

Recommended exercise 2

```
library(ISLR) # Package with data for an Introduction to Statistical
# Learning with Applications in R

# Load Credit dataset
data(Credit)

# Check column names
names(Credit)

## [1] "ID"          "Income"      "Limit"       "Rating"      "Cards"       "Age"
## [7] "Education"  "Gender"      "Student"     "Married"     "Ethnicity"   "Balance"

# Check dataset shape
dim(Credit)

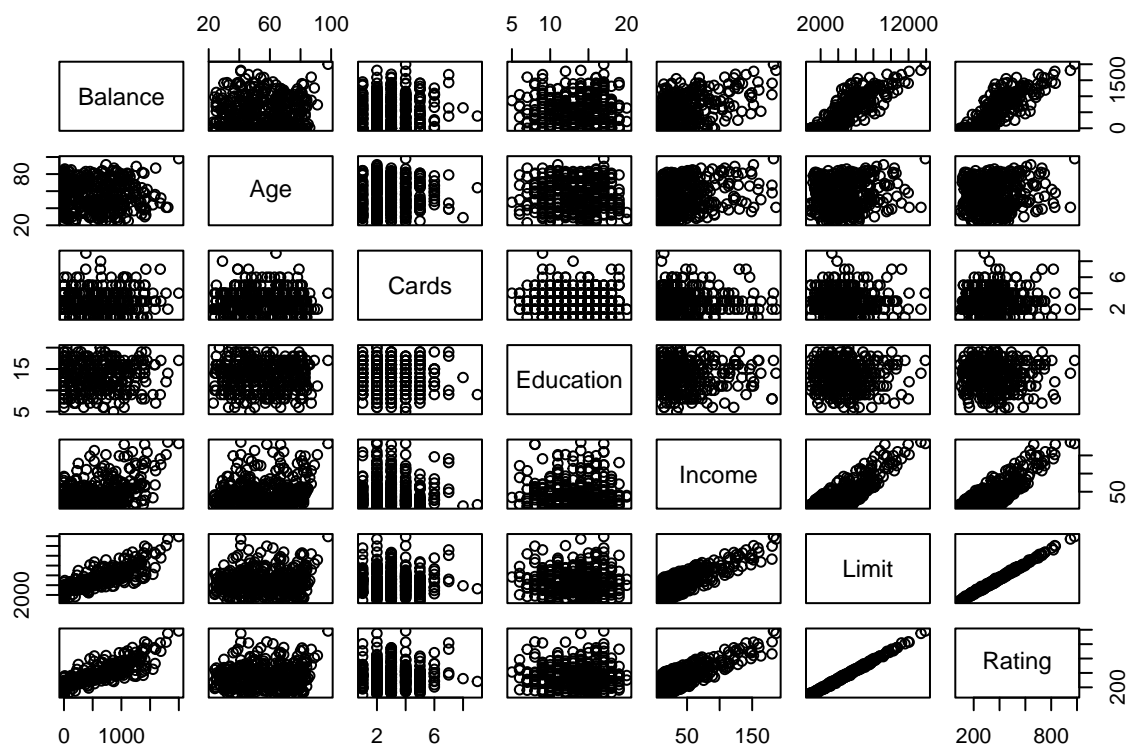
## [1] 400 12

head(Credit)

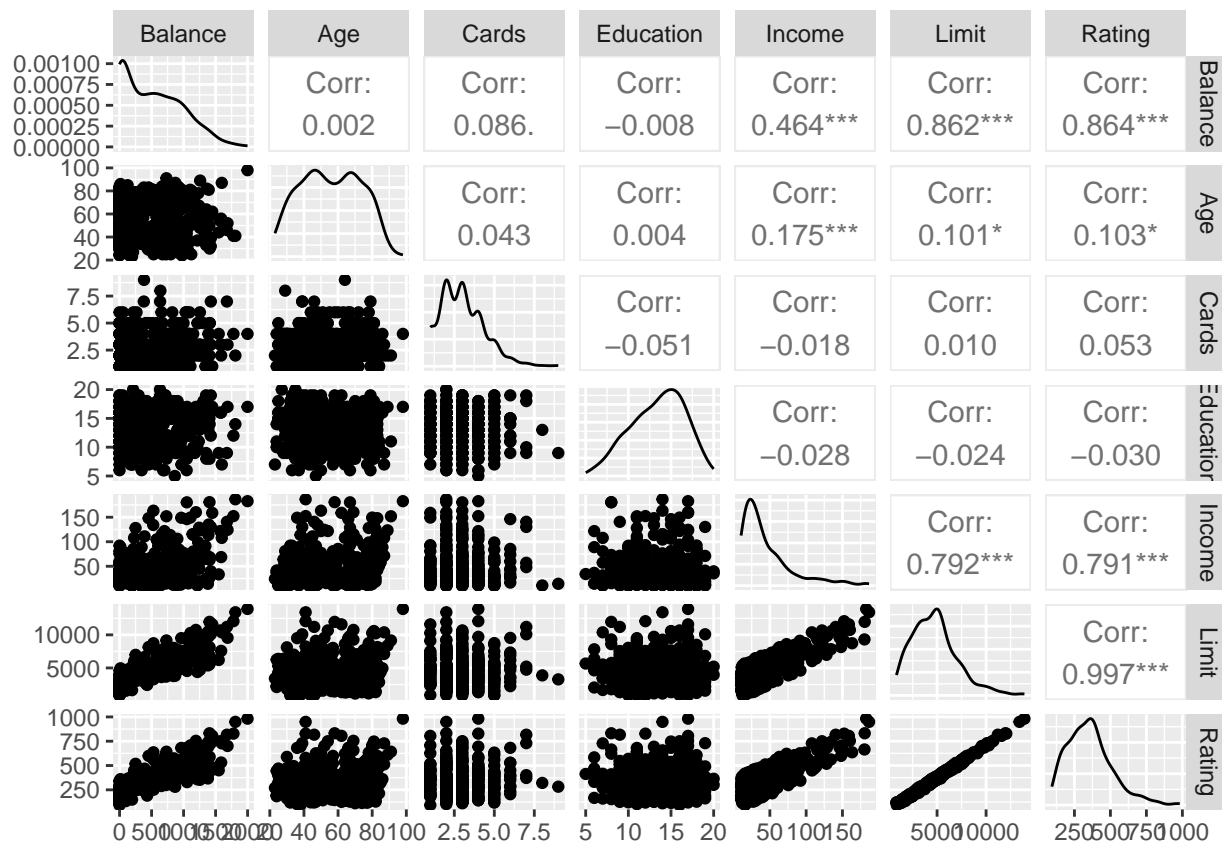
##   ID  Income Limit Rating Cards Age Education Gender Student Married Ethnicity
## 1  1  14.891  3606   283     2  34         11  Male      No      Yes  Caucasian
## 2  2 106.025  6645   483     3  82         15 Female    Yes     Yes    Asian
## 3  3 104.593  7075   514     4  71         11  Male      No      No    Asian
## 4  4 148.924  9504   681     3  36         11 Female    No      No    Asian
## 5  5  55.882  4897   357     2  68         16  Male      No     Yes  Caucasian
## 6  6  80.180  8047   569     4  77         10  Male      No     No  Caucasian
##   Balance
## 1     333
## 2     903
## 3     580
## 4     964
## 5     331
## 6    1151
```

```
# Select variable to plot
vars <- c("Balance", "Age", "Cards", "Education", "Income", "Limit", "Rating")
pairwise_scatter_data <- Credit[, vars]

# Simplest possible pairwise scatter plot
pairs(pairwise_scatter_data)
```



```
# More interesting but slower pairwise plot from package GGally
library(GGally)
ggpairs(data = pairwise_scatter_data)
```



Check here for a quick guide on getting started to ggpairs.

Recommended exercise 3

```
# Exclude 'ID' column
credit_data <- subset(Credit, select = -ID)

# Counting the dummy variables as well
credit_data_number_predictors <- 11

# Take a look at the data
head(credit_data)
```

```
##      Income Limit Rating Cards Age Education Gender Student Married Ethnicity
## 1  14.891  3606    283     2  34         11  Male      No      Yes Caucasian
## 2 106.025  6645    483     3  82         15 Female    Yes      Yes    Asian
## 3 104.593  7075    514     4  71         11  Male      No      No    Asian
## 4 148.924  9504    681     3  36         11 Female    No      No    Asian
## 5  55.882  4897    357     2  68         16  Male      No      Yes Caucasian
## 6  80.180  8047    569     4  77         10  Male      No      No    Caucasian
##      Balance
## 1         333
## 2         903
## 3         580
## 4         964
## 5         331
```

```
## 6      1151
```

```
# Summary statistics  
summary(credit_data)
```

```
##      Income      Limit      Rating      Cards  
## Min.   : 10.35   Min.    : 855   Min.    : 93.0   Min.    :1.000  
## 1st Qu.: 21.01   1st Qu.: 3088   1st Qu.:247.2   1st Qu.:2.000  
## Median : 33.12   Median : 4622   Median :344.0   Median :3.000  
## Mean   : 45.22   Mean    : 4736   Mean    :354.9   Mean    :2.958  
## 3rd Qu.: 57.47   3rd Qu.: 5873   3rd Qu.:437.2   3rd Qu.:4.000  
## Max.   :186.63   Max.    :13913   Max.    :982.0   Max.    :9.000  
##      Age      Education      Gender      Student      Married  
## Min.   :23.00   Min.    : 5.00   Male :193   No :360   No :155  
## 1st Qu.:41.75   1st Qu.:11.00   Female:207   Yes: 40   Yes:245  
## Median :56.00   Median :14.00  
## Mean   :55.67   Mean    :13.45  
## 3rd Qu.:70.00   3rd Qu.:16.00  
## Max.   :98.00   Max.    :20.00  
##      Ethnicity      Balance  
## African American: 99   Min.    : 0.00  
## Asian           :102   1st Qu.: 68.75  
## Caucasian       :199   Median : 459.50  
##                  Mean   : 520.01  
##                  3rd Qu.: 863.00  
##                  Max.   :1999.00
```

```
# Create train and test set indexes  
set.seed(1)  
train_perc <- 0.75  
credit_data_train_index <- sample(  
  nrow(credit_data),  
  round(nrow(credit_data) * train_perc)  
)
```

```
# Create train and test set  
credit_data_training <- credit_data[credit_data_train_index, ]  
credit_data_testing  <- credit_data[-credit_data_train_index, ]
```

```
library(leaps)
```

```
# Perform best subset selection using all the predictors and the training data  
best_subset_method <- regsubsets(Balance ~ .,  
  data = credit_data_training,  
  nvmax = credit_data_number_predictors  
)
```

```
# Save summary obj  
best_subset_method_summary <- summary(best_subset_method)
```

```
# Plot RSS, Adjusted R^2, C_p and BIC
```

```
par(mfrow = c(2, 2))  
plot(best_subset_method_summary$rss,  
  xlab = "Number of Variables",  
  ylab = "RSS",  
  type = "l")
```

```





)
plot(best_subset_method_summary$adjr2,
     xlab = "Number of Variables",
     ylab = "Adjusted RSq",
     type = "l"
)
bsm_best_adj2 <- which.max(best_subset_method_summary$adjr2)

points(bsm_best_adj2,
       best_subset_method_summary$adjr2[bsm_best_adj2],
       col = "red",
       cex = 2,
       pch = 20
)
plot(best_subset_method_summary$cp,
     xlab = "Number of Variables",
     ylab = "Cp",
     type = "l"
)
bsm_best_cp <- which.min(best_subset_method_summary$cp)

points(bsm_best_cp,
       best_subset_method_summary$cp[bsm_best_cp],
       col = "red",
       cex = 2,
       pch = 20
)
bsm_best_bic <- which.min(best_subset_method_summary$bic)

plot(best_subset_method_summary$bic,
     xlab = "Number of Variables",
     ylab = "BIC",
     type = "l"
)
points(bsm_best_bic,
       best_subset_method_summary$bic[bsm_best_bic],
       col = "red",
       cex = 2,
       pch = 20
)

```



```
# Create a prediction function to make predictions
# for regsubsets with id predictors included
predict.regsubsets <- function(object, newdata, id, ...) {
  form <- as.formula(object$call[[2]])
  mat <- model.matrix(form, newdata)
  coefi <- coef(object, id = id)
  xvars <- names(coefi)
  mat[, xvars] %*% coefi
}

# Create indexes to divide the data between folds
k <- 10
set.seed(1)
folds <- sample(k, nrow(credit_data_training), replace = TRUE)
cv.errors <-
  matrix(NA,
    nrow = k,
    ncol = credit_data_number_predictors,
    dimnames = list(NULL, paste(1:credit_data_number_predictors))
  )

# Perform CV
for (j in seq_len(k)) {
  best_subset_method <- regsubsets(Balance ~ .,
    data = credit_data_training[folds != j, ],
    nvmax = credit_data_number_predictors
  )
  for (i in seq_len(credit_data_number_predictors)) {
    pred <- predict(best_subset_method,
      credit_data_training[folds == j, ],

```

```


    id = i
  )
  cv.errors[j, i] <- mean((credit_data_training$Balance[folds == j] - pred)^2)
}
}

# Compute mean cv errors for each model size
mean.cv.errors <- apply(cv.errors, 2, mean)
mean.cv.errors

##          1          2          3          4          5          6          7          8
## 53308.978 27681.063 10497.276 9349.190 9468.743 9484.566 9410.272 9409.024
##          9         10         11
## 9437.443 9480.517 9496.783

# Plot the mean cv errors
par(mfrow = c(1, 1))
plot(mean.cv.errors, type = "b")

```



```

# Fit the selected model using the whole training data
# and compute test error

# models selected
number_predictors_selected <- 4

# Create info for lm call
variables <- names(coef(best_subset_method, id = number_predictors_selected))
variables <- variables[!variables %in% "(Intercept)"]
bsm_formula <- as.formula(best_subset_method$call[[2]])
bsm_design_matrix <- model.matrix(bsm_formula, credit_data_training)[, variables]
bsm_data_train <- data.frame(Balance = credit_data_training$Balance, bsm_design_matrix)

# Fit a standard linear model using only the selected

```

```

# predictors on the training data
model_best_subset_method <- lm(formula = bsm_formula, bsm_data_train)
summary(model_best_subset_method)

##
## Call:
## lm(formula = bsm_formula, data = bsm_data_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -160.26  -76.81  -11.21   48.15  350.49
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -5.216e+02  1.758e+01 -29.670  < 2e-16 ***
## Income      -7.856e+00  2.651e-01 -29.627  < 2e-16 ***
## Limit        2.706e-01  4.001e-03  67.622  < 2e-16 ***
## Cards        2.426e+01  3.981e+00   6.094 3.43e-09 ***
## StudentYes   4.196e+02  1.782e+01  23.542  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 96.14 on 295 degrees of freedom
## Multiple R-squared:  0.9575, Adjusted R-squared:  0.9569
## F-statistic: 1661 on 4 and 295 DF,  p-value: < 2.2e-16

# Make predictions on the test set
bsm_design_matrix_test <- model.matrix(bsm_formula, credit_data_testing)[, variables]
bsm_predictions <- predict(object = model_best_subset_method, newdata = as.data.frame(bsm_design_matrix_test))

# Compute test squared errors
bsm_squared_errors <- (credit_data_testing$Balance - bsm_predictions)^2
squared_errors <- data.frame(bsm_squared_errors = bsm_squared_errors)

# test MSE
mean(bsm_squared_errors)

## [1] 12243.75

```

Recommended exercise 4

Similar analysis as previous exercise, simply replace Best Subset Selection (`best_subset_method <- regsubsets(Balance ~ ., credit_data, nvmax = credit_data_number_predictors)`) by Forward Stepwise Selection (`regfit.fwd <- regsubsets(Balance ~ ., credit_data, nvmax = credit_data_number_predictors, method = "forward")`), Backward Stepwise Selection (`regfit.fwd <- regsubsets(Balance ~ ., credit_data, nvmax = credit_data_number_predictors, method = "backward")`) and Hybrid Stepwise Selection (`regfit.fwd <- regsubsets(Balance ~ ., credit_data, nvmax = credit_data_number_predictors, method = "seqrep")`)

Recommended exercise 5


```
library(glmnet) # Package Lasso and Elastic-Net Regularized
# Generalized Linear Models

x_train <- model.matrix(Balance ~ ., credit_data_training)[, -1]
y_train <- credit_data_training$Balance

x_test <- model.matrix(Balance ~ ., credit_data_testing)[, -1]
y_test <- credit_data_testing$Balance

ridge_mod <- glmnet(x_train, y_train, alpha = 0) # `alpha=0` is the ridge penalty.

set.seed(1)
cv.out <- cv.glmnet(x_train, y_train, alpha = 0)
plot(cv.out)
```



```
best_lambda_ridge <- cv.out$lambda.min
best_lambda_ridge


## [1] 40.24862

ridge_predictions <- predict(ridge_mod, s = best_lambda_ridge, newx = x_test)
ridge_square_errors <- as.numeric((ridge_predictions - y_test)^2)
squared_errors <- data.frame(ridge_square_errors = ridge_square_errors,
                             squared_errors)
```

Recommended exercise 6

```
lasso_mod <- glmnet(x_train, y_train, alpha = 1) # `alpha=1` is the lasso penalty.

set.seed(1)
cv.out <- cv.glmnet(x_train, y_train, alpha = 1)
plot(cv.out)
```



```
best_lambda_lasso <- cv.out$lambda.min
best_lambda_lasso
```

```
## [1] 1.380717
```

```
lasso_predictions <- predict(lasso_mod, s = best_lambda_lasso, newx = x_test)
lasso_square_errors <- as.numeric((lasso_predictions - y_test)^2)
squared_errors <- data.frame(lasso_square_errors = lasso_square_errors,
                             squared_errors)
```

Recommended exercise 7

```
x <- model.matrix(Balance ~ ., credit_data)[, -1]


credit_pca <- prcomp(x, center = TRUE, scale. = TRUE)

print(credit_pca)
```

```
## Standard deviations (1, ..., p=11):
## [1] 1.66007642 1.26685832 1.05356810 1.04926273 1.00322222 0.98576693
## [7] 0.97830708 0.90714714 0.63722533 0.51174012 0.04617646
##
## Rotation (n x k) = (11 x 11):
##                                PC1          PC2          PC3          PC4
```

## Income	-0.542206953	0.029036783	-0.033270648	-6.564051e-05	
## Limit	-0.586332930	0.017502630	-0.024351723	4.678929e-02	
## Rating	-0.586751867	0.014971105	-0.004630758	3.687909e-02	
## Cards	-0.019086978	-0.008549632	0.479005750	-2.720228e-01	
## Age	-0.122783390	-0.071116603	0.107188498	-4.787335e-01	
## Education	0.026797471	0.096557225	-0.475418336	1.990653e-01	
## GenderFemale	-0.002519860	0.052811098	-0.334014058	-4.207748e-02	
## StudentYes	0.002276904	0.125422970	-0.618650527	-2.963169e-01	
## MarriedYes	-0.026218561	0.094278214	0.125718135	7.389864e-01	
## EthnicityAsian	0.032769895	0.696759512	0.105703127	6.686132e-03	
## EthnicityCaucasian	-0.004070799	-0.686505857	-0.100240068	1.338718e-01	
##	PC5	PC6	PC7	PC8	PC9
## Income	-0.02816858	0.02297156	-0.04086888	0.03502243	-0.016018928
## Limit	0.02393728	0.06109959	0.02753603	-0.07998103	-0.010697575
## Rating	0.03044748	0.04901285	0.06298342	-0.07474080	-0.005366527
## Cards	0.07450235	-0.28313105	0.77070237	-0.10917776	0.005357720
## Age	-0.29468570	-0.58353604	-0.35860755	0.41270188	-0.048994454
## Education	-0.58335540	-0.40244676	0.21601791	-0.41794930	-0.021973159
## GenderFemale	0.74620452	-0.51375214	-0.10203846	-0.22746095	0.014513597
## StudentYes	0.05874438	0.20236658	0.42777847	0.53366278	0.022068488
## MarriedYes	0.04850438	-0.32419986	0.13571418	0.53676497	0.119017609
## EthnicityAsian	0.02125450	0.01284830	-0.04334986	0.01824866	-0.706522468
## EthnicityCaucasian	0.04400214	-0.02306227	0.10322555	0.06987098	-0.694731116
##	PC10	PC11			
## Income	0.836411394	0.0017092799			
## Limit	-0.379489022	0.7053633132			
## Rating	-0.373834509	-0.7081335719			
## Cards	0.059511066	0.0305564113			
## Age	-0.102540342	0.0005901693			
## Education	0.014172918	-0.0036133922			
## GenderFemale	0.027300122	0.0001327203			
## StudentYes	-0.032119354	0.0044219212			
## MarriedYes	-0.018248384	0.0051766487			
## EthnicityAsian	-0.014783578	-0.0035849536			
## EthnicityCaucasian	0.008145839	-0.0004464620			

```
plot(credit_pca, type = "l")
```



```
summary(credit_pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  1.6601 1.2669 1.0536 1.0493 1.0032 0.98577 0.97831
## Proportion of Variance 0.2505 0.1459 0.1009 0.1001 0.0915 0.08834 0.08701
## Cumulative Proportion 0.2505 0.3964 0.4973 0.5974 0.6889 0.77727 0.86427
##              PC8      PC9      PC10     PC11
## Standard deviation  0.90715 0.63723 0.51174 0.04618
## Proportion of Variance 0.07481 0.03691 0.02381 0.00019
## Cumulative Proportion 0.93908 0.97600 0.99981 1.00000
```

The first PC explains about 25% of the variability in the data. Then the second PC explains an extra 15% of the variability in the data. From the third PC until 8th PC the extra variability explained per PC varies between 7.5% to 10%, dropping to 3.6% on the 9th PCA. So I would likely use 8 PCs for the **Credit** dataset.


Recommended exercise 8

```
library(pls)

set.seed(1)

pcr_model <- pcr(Balance ~ .,
                 data = credit_data_training,
                 scale = TRUE,
                 validation = "CV")
validationplot(pcr_model, val.type = "MSEP")
```


Balance



```
pcr_predictions <- predict(pcr_model, credit_data_testing, ncomp = 10)
pcr_square_errors <- as.numeric((pcr_predictions - credit_data_testing$Balance)^2)
squared_errors <- data.frame(pcr_square_errors = pcr_square_errors, squared_errors)
mean(pcr_square_errors)
```

```
## [1] 11578.1
```

```
library(ggplot2)
library(reshape2)
ggplot(melt(squared_errors)) +
  geom_boxplot(aes(variable, value))
```




Recommended exercise 9

```
library(pls)

set.seed(1)

plsr_model <- plsr(Balance ~ .,
  data = credit_data_training,
  scale = TRUE,
  validation = "CV")
validationplot(plsr_model, val.type = "MSEP")
```


Balance



```
plsr_predictions <- predict(plsr_model, credit_data_testing, ncomp = 3)
plsr_square_errors <- as.numeric((plsr_predictions - credit_data_testing$Balance)^2)
squared_errors <- data.frame(plsr_square_errors = plsr_square_errors,
                             squared_errors)
mean(plsr_square_errors)
```

```
## [1] 12476.32
```

```
ggplot(melt(squared_errors)) +
  geom_boxplot(aes(variable, value))
```



```
colMeans(squared_errors)
```

```
##  plsr_square_errors  pcr_square_errors lasso_square_errors ridge_square_errors
##           12476.32           11578.10           12077.15           15742.83
##  bsm_squared_errors
##           12243.75
```